


Lecture 8

Convolution Neural Network

Image Classification



(All the images to be classified have the same size.)

Cross-Entropy

Intuitively Understanding the Cross Entropy

$$H(P^*|P) = - \sum_i \underbrace{P^*(i)}_{\text{TRUE CLASS DISTIRBUTION}} \log \underbrace{P(i)}_{\text{PREDICTED CLASS DISTIRBUTION}}$$

Intuitively Understanding the KL divergence

KL is not symmetric

$$D_{KL}(P||Q) = \sum_i P(i) \log \frac{P(i)}{Q(i)}$$

$$p_1 \log \frac{p_1}{q_1} + p_2 \log \frac{p_2}{q_2}$$






Image Classification




Fully Connected Network



Do we really need “fully connected” in image processing?


Observation 1




Perhaps human also identify birds in a similar way ... ☺

Observation 1

Need to see the whole image?




A neuron does not have to see the whole image.




Some patterns are much smaller than the whole image.

Simplification 1




Simplification 1

- Can different neurons have different sizes of receptive field?
- Cover only some channels?
- Not square receptive field?




Simplification 1 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).



Convolution Spatial Dimensions




Input: 7x7

Filter: 3x3

Q: How big is output?

Convolution Spatial Dimensions




Input: 7x7

Filter: 3x3

Q: How big is output?

Output: 5x5

Convolution Spatial Dimensions



Input: 7x7

Filter: 3x3

Q: How big is output?

Output: 5x5

In general:

Input: W

Filter: K



Padding: P

Stride: S


Output: $\text{floor}((W - K + 2P)/S) + 1$

Observation 2


- The same patterns appear in different regions.





Simplification 2



Simplification 2




Two neurons with the same receptive field would not share parameters.



Simplification 2 – Typical Setting


Each receptive field has a set of neurons (e.g., 64 neurons).



Simplification 2 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).


Each receptive field has the neurons with the same set of parameters.



<https://cs231n.github.io/assets/conv-demo/>


Stacking Convolution Filters

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!


Benefit of Convolutional Layer



- Some patterns are much smaller than the whole image.
- The same patterns appear in different regions.

Convolutional Layer

Another story based on filter ☺



Convolutional Layer

Consider channel = 1
(black and white image)

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

⋮ ⋮


(The values in the filters
are unknown
parameters.)

Convolutional Layer


stride=1



6 x 6 image



Filter 1



Convolutional Layer

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image


-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


Do the same process
for every filter




Convolutional Layer



Multiple Convolutional Layer



Multiple Convolutional Layer




64
filters

Convolution


Convolution

⋮
⋮

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



Receptive Fields




Convolutional Layer

<u>Neuron Version Story</u>	<u>Filter Version Story</u>
Each neuron only considers a receptive field.	There are a set of filters detecting small patterns.
The neurons with different receptive fields share the parameters.	Each filter convolves over the input image.

They are the same story.


We're still doing matrix multiplications, just localized & shared

Recall one neuron in FC layer:




w takes the
entire image!

With Conv layer:



What needs to be learned?



Observation 3

- Subsampling the pixels will not change the object



Pooling – Max Pooling

No learnable parameters!

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1


-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3


Convolutional Layers + Pooling



The whole CNNs



Application: Playing Go



Next move
(19 × 19 positions)
19 × 19 classes

48 channels in
Alpha Go

Black: 1
white: -1
none: 0


Fully-connected
network can be used

But CNN performs much
better.



Why CNN for Go playing?

- Some patterns are much smaller than the whole image

Alpha Go uses 5×5 for first layer



- The same patterns appear in different regions.



```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

for epoch in range(2): # 多批次循环

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # 获取输入
        inputs, labels = data

        # 梯度置0
        optimizer.zero_grad()

        # 正向传播, 反向传播, 优化
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

To learn more ...

- CNN is not invariant to scaling and rotation (we need data augmentation 😊).





Jürgen Schmidhuber (28 Sep 2025)
 Pronounce: You_again Shmidhoobuh
 Technical Report IDSIA-09-25, IDSIA

AI Blog
 @SchmidhuberAI
 juergen@idsia.ch

Who invented deep residual learning?

Modern AI is based on [deep artificial neural networks](#) (NNs).^[DLH] As of 2025, the most cited scientific article of the 21st century is an NN paper on *deep residual learning with residual connections*.^[MOST25,25b] Who invented this? Here is the timeline of the evolution of deep residual learning:

课堂练习题

假设一个简单的CNN用于处理MNIST手写数字数据集（输入为 28×28 像素的灰度图像，通道数为1）

网络结构如下：

1. 卷积层1：16个 3×3 卷积核，步幅（stride）为1，填充（padding）为1
2. 池化层1： 2×2 最大池化，步幅为2，无填充（padding=0）
3. 卷积层2：32个 3×3 卷积核，步幅为1，填充为1
4. 池化层2： 2×2 最大池化，步幅为2，无填充

课堂练习题

1. 求卷积层1、池化层1、卷积层2、池化层2的输出特征图尺寸（宽度、高度、通道数）
2. 计算卷积层1和卷积层2的参数数量（包括权重和偏置）
3. 若将池化层2的输出特征图展平后送入一个具有10个神经元的全连接层（对应0-9分类），求该全连接层的参数数量